

Math 140H — Honors Calculus I: Lecture Notes

Jan Reimann

2026-09-14

Table of contents

Welcome!	3
I Real Numbers	4
1 The Rational Numbers and the Number Line	5
1.1 From Counting to Measuring	5
1.2 The Number Line	6
1.3 Distances and Intervals	7
1.4 Incommensurable Lengths	8
1.5 Summary	9
2 The Real Numbers via Nested Intervals	10
2.1 Nested Intervals	10
2.2 The Nested Interval Axiom	11
2.3 Decimal Expansions	12
2.4 Other Bases	13
2.5 Summary	14
3 Jupyter Activity: p-ary Expansions	15
3.1 Before we start: how to use a Jupyter notebook	15
3.2 Importing Python modules	16
3.3 From a rational number to its p -ary expansion	17
3.4 Checking the examples	20
3.5 A few more instructive examples	20
3.6 Try it yourself	21
3.7 Food for thought	22
II Functions	23
4 The Concept of Function	24
4.1 Functions, Domain, and Range	24
4.2 Injective Functions	25
4.3 Monotone Functions	25
4.4 Even and Odd Functions	26

4.5	Building New Functions from Old	26
4.6	Summary	28
5	The Elementary Functions	29
5.1	Polynomials	29
5.2	Rational Functions	30
5.3	Algebraic and Power Functions	30
5.4	Exponential Functions	31
5.5	Logarithms	31
5.6	Trigonometric Functions	32
5.7	Summary	32
6	Continuity	34
6.1	The ε - δ Definition	35
6.2	Types of Discontinuities	36
6.3	Summary	37
7	Activity: Properties of Functions	38
7.1	Background	38
7.2	Step 1: Draw the graphs	38
7.3	Step 2: Make an educated guess	38
7.4	Verification: injectivity	39
7.5	Verification: even / odd	40
7.6	Verification: monotonicity	40
7.7	Continuity: the ε - δ game	40
7.8	Discontinuities	40
III	Limits	41
8	Sequences and Convergence	42
8.1	Sequences	42
8.2	The Limit of a Sequence	42
8.3	Convergence to a General Limit	43
8.4	Limit Laws	45
8.5	Comparison and the Squeeze Theorem	45
8.6	Monotone and Bounded Sequences	46
8.7	Applications	47
8.8	Summary	49
9	Convergence of Sequences: Reading Data	50
9.1	A Simple Query Function	53
9.2	Task 1: Explore the Data	53
9.3	Task 2: The ε Challenge	54

9.4	Dataset 2: Damped Oscillator	55
9.5	Task 3: Explore the Sequence	55
9.6	Task 4: The ε Challenge, Again	56
9.7	Task 5: Convergence of $\sqrt[n]{p}$	57

Welcome!

These notes accompany Math 140H, Honors Calculus I, Fall 2026. Each lecture here corresponds to one class session and covers more or less the material presented in class, in the same order, so that it can serve as a compact study reference. HTML “hardcopies” of the Jupyter notebook activities from class are also included. Lecture notes will generally be available immediately **after** the corresponding class session.

State College, Fall 2026

Part I

Real Numbers

1 The Rational Numbers and the Number Line

Companion reading

This lecture corresponds to C&J §1.1a, together with the neighborhood definition from §1.1d — covered a bit early here since we will use it constantly for the rest of the semester.

Calculus is built on the real numbers, so before we can talk about limits, derivatives, or integrals we need to understand exactly what a real number *is* — and, just as importantly, why the rational numbers alone are not enough. This lecture makes that precise: we recall how the natural numbers extend to the rationals, set up the number line, and then prove that a length as ordinary as the hypotenuse of a simple right triangle cannot be a rational number at all.

1.1 From Counting to Measuring

There are two major uses for numbers: **counting** and **measuring**. The natural numbers $\mathbb{N} = \{1, 2, 3, \dots\}$ are the tool for counting — abstract symbols for how many elements are in a collection, stripped of any reference to what is actually being counted. They come equipped with addition and multiplication satisfying a short list of familiar rules.

Remark 1.1. For all natural numbers a, b, c :

- $a + b = b + a$ and $ab = ba$ (commutative laws)
- $a + (b + c) = (a + b) + c$ and $a(bc) = (ab)c$ (associative laws)
- $a(b + c) = ab + ac$ (distributive law)
- $a + c = b + c \Rightarrow a = b$ (cancellation law)

These are the *rules* of arithmetic, not its *content* — everything else in this course is built on top of them, and we will not dwell on them further.

Natural numbers can be added and multiplied, but not always subtracted or divided. There is no natural number a with

$$5 + a = 2,$$

so subtraction is not always possible in $\mathbb{N}$. Extending $\mathbb{N}$ to the **integers** $\mathbb{Z} = \{\dots, -2, -1, 0, 1, 2, \dots\}$ fixes this: now $5 + (-3) = 2$.

In $\mathbb{Z}$, we still cannot always divide: there is no integer $b \in \mathbb{Z}$ with

$$5 \cdot b = 2.$$

Extending $\mathbb{Z}$ to the **rational numbers** fixes this in turn.

Definition 1.1.

A **rational number** is a number that can be written as p/q where $p, q \in \mathbb{Z}$ and $q \neq 0$. Every rational number has a unique representation p/q with $q > 0$ and $\gcd(p, q) = 1$ (no common factor larger than 1); we call this representation **in lowest terms**.

Within $\mathbb{Q}$, all four rational operations — addition, subtraction, multiplication, and division (except by 0) — can be carried out freely, and they obey exactly the same laws as in Remark 1.1. In this precise sense, $\mathbb{Q}$ is the smallest extension of the natural numbers in which counting-type arithmetic becomes measuring-type arithmetic.

1.2 The Number Line

To turn numbers into geometry, fix a line L , a point O on it called the **origin**, and a second point I to the right of O ; the segment OI serves as our unit scale for measuring — a “prototype meter.” The direction from O to I is called **positive**.

Every point on L is now uniquely determined by two pieces of information: its **distance** from O , and its **direction** (left or right) from O . For example, the point at distance $\frac{1}{2}$ in the positive direction represents the rational number $\frac{1}{2}$; the point at distance 2 in the positive direction represents 2.

Definition 1.2.

The distance of a number x from the origin is written $|x|$ and called the **absolute value** of x :

$$|x| = \begin{cases} x & x \geq 0 \\ -x & x < 0. \end{cases}$$

Note $|x| \geq 0$ always, with equality only when $x = 0$.

Every rational point can in fact be *constructed*, not just described: to locate p/q on L with ruler and compass, construct a subdivision of the segment OI into q equal parts, then extend the part closest to O , p times. We will identify each rational number with its representation on L .

1.3 Distances and Intervals

For $x, y \in L$, the number $|x - y|$ is the **distance between x and y** . Observe: if $x, y \in \mathbb{Q}$, then $|x - y| \in \mathbb{Q}$ as well — distances between rational points are themselves rational.

Definition 1.3.

For $a < b$, the **interval** between a and b comes in four flavors:

- open: $(a, b) := \{x : a < x < b\}$ (both endpoints excluded)
- closed: $[a, b] := \{x : a \leq x \leq b\}$ (both endpoints included)
- half-open: $[a, b) := \{x : a \leq x < b\}$ and $(a, b] := \{x : a < x \leq b\}$

Definition 1.4.

Given a number a and $\varepsilon > 0$, the **ε -neighborhood of a** is the open interval

$$\{y : |y - a| < \varepsilon\} = (a - \varepsilon, a + \varepsilon),$$

i.e. the set of points less than ε away from a .

We will meet ε -neighborhoods constantly once we get to limits — keep this definition close by.

1.3.1 Density

Proposition 1.1.

Between any two distinct rational numbers $a < b$ there is another rational number c with $a < c < b$. Consequently, between a and b there are infinitely many rational numbers.

Proof. Let $c = \frac{a+b}{2}$. Since $a, b \in \mathbb{Q}$ and $\mathbb{Q}$ is closed under addition and division by 2, $c \in \mathbb{Q}$; and since $a < b$, averaging gives $a < c < b$.

Now repeat the argument on the pair a, c to get a rational number strictly between a and c , hence strictly between a and b . Doing this again and again produces infinitely many rational numbers between a and b . $\square$

We say the rationals are **dense** on the number line: no matter how closely you look, you will always find another one.

1.4 Incommensurable Lengths

Given how constructible and how dense $\mathbb{Q}$ is, a natural question arises: is there anything “left in between” the rational numbers? Are there any *gaps*? Clearly, such a gap could not itself be an interval — density rules that out, since any interval of positive length contains infinitely many rational points. If a gap exists at all, it can only be a single missing point.

Remark 1.2. The discovery that not every length is a rational multiple of a chosen unit is due to the school of Pythagoras, in the fifth or sixth century BCE. Consider a right triangle with both legs of length 1; by the Pythagorean theorem, its hypotenuse a satisfies

$$1^2 + 1^2 = a^2 = 2,$$

and a can certainly be constructed with ruler and compass. The theorem below shows this length cannot be written as p/q for integers p, q — the hypotenuse and the legs of this triangle are, in the language of the ancients, **incommensurable**: there is no common unit of which both are integer multiples.

Theorem 1.1.

There is no rational number x with $x^2 = 2$. Equivalently, $\sqrt{2} \notin \mathbb{Q}$.

Proof. Suppose, toward a contradiction, that $\sqrt{2} = p/q$ for integers p, q with $q \neq 0$ and $\gcd(p, q) = 1$ (i.e. p/q is in lowest terms).

1. Square both sides and clear the denominator: $p^2 = 2q^2$.
2. Conclude that p^2 is even — if p were odd, p^2 would be odd, but $p^2 = 2q^2$ is even. Write $p = 2r$ for some integer r .
3. Substitute $p = 2r$ into the equation from Step 1 and simplify to get an equation for q^2 in terms of r : $(2r)^2 = 2q^2 \Rightarrow 4r^2 = 2q^2 \Rightarrow q^2 = 2r^2$.
4. Conclude that q is also even, by the same reasoning as in Step 2.
5. Steps 2 and 4 show that 2 divides both p and q . This contradicts $\gcd(p, q) = 1$, which we assumed when we wrote p/q in lowest terms.
6. Conclude: $\sqrt{2}$ is irrational — no integers p, q satisfy $\sqrt{2} = p/q$, so $\sqrt{2} \notin \mathbb{Q}$.

□

Exercise 1.1. Adapt the proof of Theorem 1.1 to show that $\sqrt{3}$ is irrational. Then try to see where the argument would break if you attempted the same proof for $\sqrt{4}$ — since $\sqrt{4} = 2$ is certainly rational, something in the argument *must* fail. Where, exactly?

1.5 Summary

- The rationals $\mathbb{Q}$ satisfy all the usual laws of arithmetic (Remark 1.1) and are **dense**: between any two rationals lies another (Proposition 1.1).
- Distances $|x - y|$, intervals, and ε -neighborhoods (Definition 1.3, Definition 1.4) give us the basic vocabulary for talking about “closeness” on the number line — vocabulary we will lean on heavily once limits arrive.
- Density notwithstanding, $\mathbb{Q}$ does not exhaust the number line: Theorem 1.1 exhibits an ordinary, constructible length that is not a rational number.
- The rationals therefore leave *gaps* in the line. Lecture 2 closes these gaps rigorously, using **nested intervals** to construct the real numbers.

2 The Real Numbers via Nested Intervals

Companion reading

This lecture corresponds to C&J §1.1b–c.

Last time we found a gap: the length $\sqrt{2}$ is constructible with ruler and compass, yet Theorem 1.1 shows it is not a rational number. So what *is* $\sqrt{2}$, as a number? We are used to describing such quantities through infinite decimal expansions, e.g.

$$\pi = 3.1415926 \dots,$$

but so far this is just a string of digits — it is not clear what it *means* mathematically. The key idea of this lecture is that an infinite decimal (or any other description of an irrational number) is really a recipe for **approximating** a point by rational numbers, to any desired degree of accuracy. We make this precise using shrinking intervals with rational endpoints.

2.1 Nested Intervals

Suppose x is a point on the number line that we want to pin down using only rational numbers. Start with any closed interval $I_1 = [a_1, b_1]$ with $a_1, b_1 \in \mathbb{Q}$ and $a_1 \leq x \leq b_1$. Now choose a subinterval $I_2 = [a_2, b_2] \subseteq I_1$ that still contains x :

$$a_1 \leq a_2 \leq x \leq b_2 \leq b_1,$$

for instance by splitting I_1 in half and keeping whichever half contains x . Repeating this forever produces a chain

$$I_1 \supseteq I_2 \supseteq I_3 \supseteq \dots, \quad x \in I_n \text{ for every } n.$$

For this to genuinely *pin down* x — rather than just hover near it — we need one more ingredient: the intervals must shrink all the way down to a point.

Definition 2.1.

A **nested sequence of closed intervals with rational endpoints** is a sequence $I_n = [a_n, b_n]$, $n = 1, 2, 3, \dots$, with $a_n, b_n \in \mathbb{Q}$, such that

1. $I_1 \supseteq I_2 \supseteq I_3 \supseteq \dots$, and
2. the lengths $|I_n| := b_n - a_n$ tend to 0 as $n \rightarrow \infty$: for every $\varepsilon > 0$ there is an N such that $|I_n| < \varepsilon$ for all $n \geq N$.

Proposition 2.1.

If a nested sequence of closed intervals (I_n) contains a point in every I_n , that point is unique: there is at most one x with $x \in I_n$ for all n .

Proof. Suppose x and y both lie in every I_n , and toward a contradiction suppose $x \neq y$. Let $\varepsilon_0 = |x - y| > 0$. Since $|I_n| \rightarrow 0$, there is some n with $|I_n| < \varepsilon_0$. But $x, y \in I_n$ means $|x - y| \leq |I_n| < \varepsilon_0 = |x - y|$, which is absurd. Hence $x = y$. $\square$

So a nested sequence of intervals can “address” at most one point — and, geometrically, every point x on the line does arise this way:

Theorem 2.1.

Every real number x is the unique point contained in every interval of some nested sequence of closed intervals with rational endpoints.

Proof. Uniqueness is Proposition 2.1. For existence, pick any rational bracket $I_1 = [a_1, b_1]$ with $a_1 \leq x \leq b_1$, and repeatedly bisect: given $I_n = [a_n, b_n]$, let $m_n = \frac{a_n + b_n}{2}$ (rational, since $a_n, b_n \in \mathbb{Q}$), and set

$$I_{n+1} = \begin{cases} [a_n, m_n] & \text{if } x \leq m_n \\ [m_n, b_n] & \text{if } x > m_n. \end{cases}$$

Then $x \in I_{n+1} \subseteq I_n$ and $|I_{n+1}| = |I_n|/2$, so $|I_n| = |I_1|/2^{n-1} \rightarrow 0$. This is a nested sequence of closed intervals with rational endpoints containing x . $\square$

In this sense, every real number can be described — completely and unambiguously — by an infinite sequence of rational numbers (the endpoints a_n, b_n).

2.2 The Nested Interval Axiom

The construction above starts from a real number x and produces a nested sequence around it. What about the converse: if we are simply *handed* a nested sequence of closed intervals with rational endpoints, must there be a point in $\bigcap_n I_n$ at all?

For **open** intervals, the answer is no: $I_n = (0, \frac{1}{n})$ satisfies $I_1 \supseteq I_2 \supseteq \dots$ and $|I_n| = \frac{1}{n} \rightarrow 0$, yet no point lies in every I_n — any candidate $x > 0$ fails once $n > 1/x$, and $x = 0$ is excluded from every I_n by openness.

For **closed** intervals, this is far less obvious, and cannot be deduced from the algebraic properties of $\mathbb{Q}$ alone (after all, $\mathbb{Q}$ itself has a gap at $\sqrt{2}$). We take it as a basic postulate about the number line.

Axiom (Nested Interval Property)

If $I_1 \supseteq I_2 \supseteq I_3 \supseteq \dots$ is a nested sequence of closed intervals with rational endpoints, then

$$\bigcap_{n \in \mathbb{N}} I_n \neq \emptyset,$$

i.e. there is a point x contained in every I_n .

This is a genuine **axiom**: a basic postulate we accept without proof, rather than a consequence of things we already know. Together with Proposition 2.1, it guarantees that no gaps exist on the number line — every nested sequence of rational brackets zeroes in on exactly one real number, no more and no less. This completeness of the number line is fundamental for everything that follows in this course; we will meet it again, in other guises, once we get to limits.

2.3 Decimal Expansions

Decimal expansions are simply a very concrete, systematic way of producing a nested sequence of intervals — one that subdivides by tens at every step.

Definition 2.2.

Let x be a point on the number line. Set

$$c_0 := \lfloor x \rfloor,$$

the greatest integer with $c_0 \leq x \leq c_0 + 1$, and $I_0 = [c_0, c_0 + 1]$, so $|I_0| = 1$.

Divide I_0 into ten equal parts using the points $c_0 + \frac{1}{10}, c_0 + \frac{2}{10}, \dots, c_0 + \frac{9}{10}$, and let $c_1 \in \{0, 1, \dots, 9\}$ be the digit with

$$c_0 + \frac{c_1}{10} \leq x \leq c_0 + \frac{c_1}{10} + \frac{1}{10},$$

and set I_1 to be this interval, so $|I_1| = \frac{1}{10}$.

Iterating, having chosen digits $c_1, \dots, c_{n-1} \in \{0, \dots, 9\}$, subdivide I_{n-1} into ten equal parts and let $c_n \in \{0, \dots, 9\}$ be the digit with

$$c_0 + \frac{c_1}{10} + \dots + \frac{c_n}{10^n} \leq x \leq c_0 + \frac{c_1}{10} + \dots + \frac{c_n}{10^n} + \frac{1}{10^n},$$

and let I_n be this interval, so $|I_n| = 10^{-n}$.

By construction $(I_n)_{n \geq 0}$ is exactly a nested sequence of closed intervals with rational endpoints in the sense of Definition 2.1: $I_0 \supseteq I_1 \supseteq I_2 \supseteq \dots$, each contains x , and $|I_n| = 10^{-n} \rightarrow 0$. We abbreviate this whole sequence by writing

$$x = c_0 + 0.c_1c_2c_3 \dots,$$

the **decimal expansion** of x — a notational shortcut for the nested intervals, not a new kind of number. (Careful when $c_0 < 0$: the digits after the point are still *added* to c_0 . For instance $x = -2.75$ has $c_0 = \lfloor -2.75 \rfloor = -3$ and $x = -3 + 0.25$, not $-3 - 0.25$.)

Combining Definition 2.2 with Proposition 2.1 and the Nested Interval Axiom:

- Every real number has an (infinite) decimal expansion.
- Conversely, every such expansion — i.e. every choice of $c_0 \in \mathbb{Z}$ and digits $c_1, c_2, \dots \in \{0, \dots, 9\}$ — is a nested sequence of closed intervals with rational endpoints, so by the axiom it defines a real number.

Decimal expansions are not always unique

Some real numbers have *two* decimal expansions: for instance

$$1.000 \dots = 0.999 \dots$$

Exercise 2.1. Why does $1.000 \dots = 0.999 \dots$? (Hint: what would $1 - 0.999 \dots$ have to be, if it were positive?) More generally, can you characterize *exactly* which real numbers have two decimal expansions?

Solution. If $1 - 0.999 \dots = \delta > 0$, then δ would have to be smaller than 10^{-n} for every n (since $0.\underbrace{9 \dots 9}_n = 1 - 10^{-n}$ already exceeds $1 - \delta$), forcing $\delta = 0$ — contradiction. So the two expansions name the same point.

In general, a real number has two decimal expansions exactly when it is a **terminating** decimal, i.e. of the form $m/10^n$ for some $m \in \mathbb{Z}$, $n \geq 0$: one expansion ends in all 0's, the other in all 9's. Every other real number has a unique decimal expansion.

2.4 Other Bases

Nothing in Definition 2.2 is special about the number 10 — we could subdivide each interval into any number $p > 1$ of equal parts instead of ten, producing digits $c_i \in \{0, 1, \dots, p-1\}$ and a **base- p (or p -ary) expansion** of x .

The case $p = 2$, the **binary expansion**, is exactly the bisection procedure from the proof of Theorem 2.1: at each step we only need to decide whether x lies in the lower or upper half of the current interval.

Example 2.1.

Let's find the first three binary digits of $x = 0.76$. Start with $I_0 = [0, 1]$, so $c_0 = 0$.

- Midpoint 0.5: since $0.76 \geq 0.5$, digit $c_1 = 1$ and $I_1 = [0.5, 1]$.
- Midpoint 0.75: since $0.76 \geq 0.75$, digit $c_2 = 1$ and $I_2 = [0.75, 1]$.
- Midpoint 0.875: since $0.76 < 0.875$, digit $c_3 = 0$ and $I_3 = [0.75, 0.875]$.

So $0.76 = 0.110\dots$ in base 2.

Exercise 2.2. Find the first three fractional digits of the base-2 expansion of $x = 0.3$, and the first three fractional digits of the base-3 expansion of $x = 0.76$.

2.5 Summary

- A **nested sequence of closed intervals with rational endpoints** (Definition 2.1) can contain at most one point (Proposition 2.1), and every real number arises as the unique point of some such sequence (Theorem 2.1).
- The converse — that every nested sequence of closed intervals with rational endpoints *does* contain a point — is not provable from what came before; it is the **Nested Interval Axiom**, the basic postulate that guarantees there are no gaps in the number line.
- **Decimal expansions** (Definition 2.2) are the special case of base-10 subdivision; every real number has one, though numbers of the form $m/10^n$ have exactly two (Exercise 2.1).
- The same construction works in any base $p > 1$, giving p -ary expansions; base 2 is exactly the bisection algorithm behind Theorem 2.1.

3 Jupyter Activity: p-ary Expansions

Math 140H, Honors Calculus I

3.1 Before we start: how to use a Jupyter notebook

This is the first notebook of the semester, so we start a quick orientation:

- A notebook is a sequence of **cells**. This text lives in a *markdown cell* (formatted text); cells with **code** in the toolbar are *code cells* (Python).
- Click a cell and press **Shift+Enter** to run it and move to the next one. Code cells must be run in order the first time through, top to bottom — a later cell may use a function or variable that an earlier cell defined.
- If something breaks, it is often because a cell was skipped or run twice out of order. Use the menu command *Run All* (or *Restart Kernel and Run All*) to run the whole notebook cleanly from the top.
- You do **not** need to know Python already. Every function below is commented line by line — read the comments (the lines starting with **#**) as a description of the math, and the code as one particular way of writing it down.

3.1.1 Code cells

You can think of a code cell very much like a fancy calculator. Move to the next cell and run it.

```
16+(53*8.7777-128)
```

```
353.2181
```

Careful: In Python, exponentiation is ******, not **^**.

```
2**4
```

```
16
```

```
2^4
```

6

`^` in Python is the operation called **bitwise XOR**; that explains the strange result above.

We will learn more about Jupyter notebooks (and a little bit more about Python) in the course of this semester. Today, we will simply experiment a bit with pre-defined Python function that converts rational numbers to p -ary expansions. Go ahead and run the cells below one after the other (do not forget to read the text inbetween.)

3.2 Importing Python modules

A lot functionality in Python is provided by *modules*. The following cell loads the modules we need for this notebook.

```
# We import three tools before doing anything else:
#
#   math      -- the standard library module with math.floor(y), the greatest integer <=
#               (We need floor of *exact fractions*, not of floating-point approximation
#               see the note below.)
#   Fraction  -- from the 'fractions' module. A Fraction stores a rational number as an e
#               numerator/denominator pair, e.g. Fraction(1, 3) is *exactly* one third,
#               This matters a lot here: a computer's ordinary decimal numbers (called '
#               are themselves binary approximations, so 1/3 as a float is already sligh
#               and the tiny error would snowball after enough steps of the algorithm be
#               Using Fraction means every digit we compute is provably correct, not jus
#   numpy     -- imported as 'np' by convention. We use it only to store the list of digi
#               compute as a numpy array, which is a convenient container for a sequence
import math
from fractions import Fraction
import numpy as np
```

The cells below show the advantage of the `Fraction` function:

```
1/2+1/3 # will use decimal expansions
```

```
0.8333333333333333
```

```
Fraction(1,2)+Fraction(1,3) # will preserve the precise fraction representation
```

```
Fraction(5, 6)
```

3.3 From a rational number to its p -ary expansion

We generalize the nested interval construction to base p : given x , set

$$c_0 := \lfloor x \rfloor,$$

and then repeatedly zoom in. Having peeled off $c_0, c_1, \dots, c_{n-1}$, look at what is left over,

$$r_{n-1} = x - c_0 - \frac{1}{p}c_1 - \dots - \frac{1}{p^{n-1}}c_{n-1} \in [0, p^{-(n-1)}),$$

multiply it by p^{n-1} so it lands back in $[0, 1)$, and read off the next digit as the integer part of p times that:

$$c_n = \lfloor p \cdot (p^{n-1}r_{n-1}) \rfloor.$$

(Multiplying by p here means zooming in by a factor of p .)

The function below does exactly this, but it tracks the *rescaled* remainder (called **remainder** in the code, always kept in $[0, 1)$) rather than r_{n-1} itself, since that is what you actually multiply by p at each step:

1. **remainder** = **x** - **floor(x)**, the fractional part of x , in $[0, 1)$.
2. Repeat: multiply **remainder** by p . Its integer part is the next digit; subtract that integer part off, leaving a new **remainder** back in $[0, 1)$, and go again.

```

def rational_to_p_ary(x, p, max_digits=50):
    """
    Compute the base-p expansion of a rational number x.

    Parameters
    -----
    x : Fraction, int, or a string like "7/12"
        The rational number to expand. (Do NOT pass a plain float such as 0.76 -- floats
        already binary approximations, so we would be expanding the wrong number. Use
        Fraction(76, 100) instead. See the discussion above.)
    p : int
        The base, an integer > 1 (p = 2 for binary, p = 10 for the usual decimal expansion)
    max_digits : int, optional (default = 50)
        how many fractional digits to compute

    Returns
    -----
    integer_part : int
        c_0, the integer part of x (i.e. floor(x)).
    digits : numpy array of int
        The fractional digits c_1, c_2, c_3, ... that were computed, in order.
    repeat_start : int or None
        If the expansion is eventually periodic, the index (0-based, into `digits`) where
        repeating block begins. If the expansion terminates exactly (all later digits are 0)
        this is None.
    """
    if not isinstance(p, int) or p <= 1:
        raise ValueError("the base p must be an integer greater than 1")

    # Convert x to an exact fraction, no matter what form it was given in.
    x = Fraction(x)

    # Step 1: peel off the integer part c_0 = floor(x). What remains, x - c_0, is the
    # 'fractional part' of x and always lies in [0, 1).
    integer_part = math.floor(x)
    remainder = x - integer_part

    digits = []          # the digits c_1, c_2, ... found so far, in order
    seen_remainders = {} # remembers, for each remainder value we've seen, which step
                        # first showed up at -- this is how we notice a repeat

    repeat_start = None

```

```

for step in range(max_digits):
    # Zooming into the next sub-interval of
    # length 1/p means multiplying the remainder by p; the integer part of the result
    # exactly which of the p equal pieces x fell into, i.e. the next digit.
    remainder = remainder * p
    next_digit = math.floor(remainder)
    digits.append(next_digit)

    # Subtract off the digit we just read, leaving a fresh remainder back in [0, 1),
    # ready for the next iteration.
    remainder = remainder - next_digit

return integer_part, np.array(digits, dtype=int), repeat_start

```

The cell below defines a small helper function to print an expansion in a readable form. We deliberately write it as $c_0 + 0.c_1c_2c_3\dots$ with an explicit $+$, matching the lecture notes' warning: if c_0 is negative, the fractional digits are still *added* to it, not subtracted, so -3.25 written the ordinary way would be dangerously ambiguous. Writing $-3 + 0.25$ instead leaves no room for that mistake.

```

def format_expansion(integer_part, digits, p, repeat_start=None):
    """
    Turn the output of rational_to_p_ary into a readable string, e.g.
        format_expansion(0, np.array([1, 1, 0]), 2, None) -> '0 + 0.110 (base 2)'
        format_expansion(0, np.array([3]), 10, 0)          -> '0 + 0.(3) (base 10)' #
    """
    # Turn each digit into a one-character string. In any base p <= 10 these are ordinary
    # for p > 10 we would need letters (like hexadecimal's a-f), which we don't handle here
    digit_chars = [str(int(d)) for d in digits]

    if repeat_start is None:
        # No repeating block: just list every digit we have.
        fractional_part = "".join(digit_chars) if len(digit_chars) > 0 else "0"
    else:
        # Split into the digits before the repeat (head) and the repeating block itself
        # and wrap the tail in parentheses to mark it as repeating forever.
        head = "".join(digit_chars[:repeat_start])
        tail = "".join(digit_chars[repeat_start:])
        fractional_part = f"{head}({tail})" if head else f"({tail})"

    return f"{integer_part} + 0.{fractional_part} (base {p})"

```

3.4 Checking the examples

In the lecture notes, we found the first three *binary* digits of $x = 0.76$ by hand: $c_1 = 1$, $c_2 = 1$, $c_3 = 0$, i.e. $0.76 = 0.110\dots$ in base 2. Let's confirm that.

```
integer_part, digits, repeat_start = rational_to_p_ary(Fraction(76, 100), 2, max_digits=20)
print(format_expansion(integer_part, digits, 2, repeat_start))
print("first three fractional digits:", digits[:3], " -- matches c1, c2, c3 = 1, 1, 0 from Example 2.6")
```

```
0 + 0.110000101000111101011100001010 (base 2)
first three fractional digits: [1 1 0] -- matches c1, c2, c3 = 1, 1, 0 from Example 2.6
```

Now find the first three fractional digits of the base-2 expansion of $x = 0.3$, and of the base-3 expansion of $x = 0.76$. Work those out on paper first – then run the cell below to check yourself.

```
ip, d, r = rational_to_p_ary(Fraction(3, 10), 2, max_digits=20)
print("0.3 in base 2:", format_expansion(ip, d, 2, r))
print("first three digits:", d[:3])
print()
ip, d, r = rational_to_p_ary(Fraction(76, 100), 3, max_digits=20)
print("0.76 in base 3:", format_expansion(ip, d, 3, r))
print("first three digits:", d[:3])
```

```
0.3 in base 2: 0 + 0.01001100110011001100 (base 2)
first three digits: [0 1 0]
```

```
0.76 in base 3: 0 + 0.20211200100201102212 (base 3)
first three digits: [2 0 2]
```

3.5 A few more instructive examples

- $1/3$ in base 10 has *no* terminating part at all – the repeating block starts immediately.
- $1/6$ in base 10 has one non-repeating digit followed by a repeating block (matches $1/6 = 0.1\bar{6}$).
- $1/10$ terminates in base 10 (as expected, $1/10 = 0.1$ exactly) but is *purely repeating* in base 2 – this is the mathematical reason 0.1 cannot be stored exactly as an ordinary floating-point number, and is why we insisted on `Fraction` above instead of Python floats.

- Negative x : $x = -11/4 = -2.75$. Its integer part is $c_0 = \lfloor -2.75 \rfloor = -3$, and the digits describe 0.25, which gets *added* to -3 (giving back -2.75), not subtracted.

```
for label, value, base, n in [
    ("1/3", Fraction(1, 3), 10, 10),
    ("1/6", Fraction(1, 6), 10, 10),
    ("1/10", Fraction(1, 10), 10, 10),
    ("1/10", Fraction(1, 10), 2, 30),
    ("-11/4", Fraction(-11, 4), 10, 10),
]:
    ip, d, r = rational_to_p_ary(value, base, max_digits=n)
    print(f"{label:>6} = {format_expansion(ip, d, base, r)}")
```

```
1/3 = 0 + 0.3333333333 (base 10)
1/6 = 0 + 0.1666666666 (base 10)
1/10 = 0 + 0.1000000000 (base 10)
1/10 = 0 + 0.000110011001100110011001100110 (base 2)
-11/4 = -3 + 0.2500000000 (base 10)
```

3.6 Try it yourself

Pick your own rational number x and a base $p > 1$. Use `rational_to_p_ary` to get its p -ary expansion.

```
# Replace the values below with your own choice of x (as a Fraction), p.
x = Fraction(1, 2)
p = 10

ip, d, r = rational_to_p_ary(x, p, max_digits=30)
print(f"x in base {p}:", format_expansion(ip, d, p, r))
```

```
x in base 10: 0 + 0.500000000000000000000000000000 (base 10)
```

3.7 Food for thought

Question 1: For a fixed base p , which fractions will have a p -ary expansion that will eventually be identically 0? (such as $1/2$ in base 10) Try to formulate a hypothesis and test it using the cell above.

Question 2: Is it true that all p -ary expansions of rational numbers eventually become periodic? If yes, why is this the case?

Part II

Functions

4 The Concept of Function

With the number line in hand, we can start doing calculus — but calculus is not really about individual numbers. It is about *dependence*: how one quantity changes as another one does. The area of a circle depends on its radius, the height of a falling stone depends on the elapsed time, the pressure of a gas depends on its volume. The mathematical object that captures such a dependence is a **function**, and getting comfortable with the vocabulary now will pay off for the rest of the course.

4.1 Functions, Domain, and Range

Definition 4.1.

A **function** f is a rule that assigns to every element x of a set D — the **domain** of f — one single, unambiguously determined object $f(x)$, its **value** at x , lying in a set E — the **codomain**. We write

$$f: D \rightarrow E, \quad x \mapsto f(x).$$

In this course the domain is almost always a subset of the number line, $D \subseteq \mathbb{R}$, and the codomain is $E = \mathbb{R}$; we then call f a **real-valued function of a real variable** and write simply $f: D \rightarrow \mathbb{R}$.

The word *unambiguously* is the heart of the definition: for each input $x \in D$ there must be **exactly one** output $f(x)$ — no more (that would not be a rule) and no fewer (every point of the domain must get a value). The domain is part of the data of a function: $x \mapsto x^2$ on $\mathbb{R}$ and $x \mapsto x^2$ on $[0, \infty)$ are, strictly speaking, different functions.

Remark 4.1. When a function is specified by a formula with no domain stated, we adopt the convention that its domain is the **natural domain**: the largest set of real numbers for which the formula makes sense. For $x \mapsto 1/(x-1)$ this is $\mathbb{R} \setminus \{1\}$; for $x \mapsto \sqrt{x}$ it is $[0, \infty)$. We are always free to *restrict* to a smaller domain when we want to.

Definition 4.2.

The **range** (or **image**) of $f: D \rightarrow \mathbb{R}$ is the set of values it actually attains:

$$\text{ran}(f) = \{y \in \mathbb{R} : y = f(x) \text{ for some } x \in D\}.$$

Remark 4.2. The range need not fill up the codomain. For $f: \mathbb{R} \rightarrow \mathbb{R}$, $f(x) = x^2$, the codomain is $\mathbb{R}$ but

$$\text{ran}(f) = [0, \infty),$$

since squares are never negative and every $y \geq 0$ is $(\sqrt{y})^2$. Saying “ f maps into $\mathbb{R}$ ” is a statement about where the values are *allowed* to live; the range records where they *actually* land.

4.2 Injective Functions

Definition 4.3.

A function $f: D \rightarrow \mathbb{R}$ is **injective** (or **one-to-one**) if it sends distinct inputs to distinct outputs:

$$x \neq y \quad \Rightarrow \quad f(x) \neq f(y),$$

or equivalently, taking the contrapositive, $f(x) = f(y) \quad \Rightarrow \quad x = y$.

Example 4.1.

- $f(x) = x^3$ is injective on $\mathbb{R}$: if $x^3 = y^3$ then $x = y$ (taking cube roots).
- $g(x) = x^2$ is **not** injective on $\mathbb{R}$: $g(-1) = 1 = g(1)$, yet $-1 \neq 1$. Restricted to $[0, \infty)$, however, g is injective.

Injectivity is exactly the property that lets us run a function backwards; we return to this when we discuss inverse functions below.

4.3 Monotone Functions

Definition 4.4.

Let $f: D \rightarrow \mathbb{R}$ with $D \subseteq \mathbb{R}$. For all $x, w \in D$, we say f is

- **increasing** if $x \leq w \quad \Rightarrow \quad f(x) \leq f(w)$;
- **strictly increasing** if $x < w \quad \Rightarrow \quad f(x) < f(w)$;
- **decreasing** if $x \leq w \quad \Rightarrow \quad f(x) \geq f(w)$;
- **strictly decreasing** if $x < w \quad \Rightarrow \quad f(x) > f(w)$.

A function of any of these four types is called **monotone**.

Example 4.2.

- $f(x) = x^3$ is strictly increasing on all of $\mathbb{R}$.

- $g: [0, \pi] \rightarrow \mathbb{R}$, $g(x) = \cos x$, is strictly decreasing: as the angle runs from 0 to π , the x -coordinate on the unit circle drops steadily from 1 to -1 .

Exercise 4.1. Show that a strictly increasing function is injective (and likewise for strictly decreasing). Is the converse true — must an injective function be monotone?

4.4 Even and Odd Functions

Definition 4.5.

Suppose the domain D is symmetric about 0, i.e. $x \in D \Rightarrow -x \in D$. A function $f: D \rightarrow \mathbb{R}$ is

- **even** if $f(-x) = f(x)$ for all $x \in D$ — its graph is symmetric about the y -axis;
- **odd** if $f(-x) = -f(x)$ for all $x \in D$ — its graph is symmetric about the origin.

Example 4.3.

x^2 , $|x|$, and $\cos x$ are even; x , x^3 , and $\sin x$ are odd; $x + x^2$ is neither. Most functions are neither even nor odd.

Exercise 4.2. Show that every function f on a symmetric domain can be written *uniquely* as a sum $f = f_e + f_o$ of an even function f_e and an odd function f_o . (Hint: what must $\frac{1}{2}(f(x) + f(-x))$ and $\frac{1}{2}(f(x) - f(-x))$ be?)

4.5 Building New Functions from Old

Most functions we work with are assembled from simpler ones by a few standard operations.

4.5.1 Arithmetic combinations

Definition 4.6.

Let $f, g: D \rightarrow \mathbb{R}$ be two functions on the same domain, and $c \in \mathbb{R}$. Then

$$f + g, \quad cf, \quad f \cdot g$$

are again functions on D , defined **pointwise**: $(f + g)(x) = f(x) + g(x)$, $(cf)(x) = cf(x)$, $(f \cdot g)(x) = f(x)g(x)$. The **quotient** f/g , given by $(f/g)(x) = f(x)/g(x)$, is a function on the smaller domain

$$D \setminus \{x \in D : g(x) = 0\}.$$

4.5.2 Composition

Definition 4.7.

Let $g: D \rightarrow \mathbb{R}$ and $f: D' \rightarrow \mathbb{R}$ with $\text{ran}(g) \subseteq D'$. The **composition** $f \circ g$ (“first g , then f ”) is

$$(f \circ g)(x) := f(g(x)), \quad x \in D,$$

a function with domain $\text{dom}(f \circ g) = \text{dom}(g)$.

Example 4.4.

With $f(x) = \sqrt{x}$ and $g(x) = 1 + x^2$:

$$(f \circ g)(x) = \sqrt{1 + x^2} \quad (\text{all } x \in \mathbb{R}), \quad (g \circ f)(x) = 1 + x \quad (x \geq 0).$$

The order matters: $f \circ g$ and $g \circ f$ are different functions.

4.5.3 Inverse functions

Definition 4.8.

If $f: D \rightarrow \mathbb{R}$ is injective, then for each $y \in \text{ran}(f)$ there is **exactly one** $x \in D$ with $f(x) = y$. Sending each such y to that x defines the **inverse function**

$$f^{-1}: \text{ran}(f) \rightarrow D, \quad f^{-1}(f(x)) = x \text{ for all } x \in D, \quad f(f^{-1}(y)) = y \text{ for all } y \in \text{ran}(f).$$

Its domain is $\text{ran}(f)$ and its range is D .

Notation trap

f^{-1} denotes the **inverse function**, not the reciprocal $1/f$. For $f(x) = x^3$ we have $f^{-1}(y) = \sqrt[3]{y}$, whereas $\frac{1}{f(x)} = x^{-3}$.

Remark 4.3. The point (a, b) lies on the graph of f exactly when (b, a) lies on the graph of f^{-1} . So the graph of f^{-1} is the mirror image of the graph of f across the line $y = x$.

Example 4.5.

$f(x) = x^3$ is injective on $\mathbb{R}$ with $\text{ran}(f) = \mathbb{R}$, and $f^{-1}(y) = \sqrt[3]{y}$. By contrast $g(x) = x^2$ on $\mathbb{R}$ has no inverse (it is not injective); restricted to $[0, \infty)$ it does, namely $g^{-1}(y) = \sqrt{y}$.

4.6 Summary

- A **function** $f: D \rightarrow \mathbb{R}$ (Definition 4.1) assigns to each point of its **domain** D exactly one real value; the set of values attained is its **range** (Definition 4.2), which may be a proper subset of the codomain.
- f is **injective** (Definition 4.3) if distinct inputs give distinct outputs, and **monotone**
 - (1) if it consistently preserves or reverses order; strict monotonicity implies injectivity (Exercise 4.1).
- On a symmetric domain, f may be **even** or **odd** (Definition 4.5), i.e. symmetric about the y -axis or the origin.
- New functions are built from old ones by pointwise **arithmetic** (Definition 4.6), **composition** (Definition 4.7), and — when f is injective — by forming the **inverse** f^{-1} (Definition 4.8), whose graph is the reflection of that of f across $y = x$.

5 The Elementary Functions

Companion reading

This lecture corresponds to C&J §1.3.

Almost every function met in a first calculus course is built out of a short catalogue of standard functions: polynomials, rational functions, powers, exponentials, logarithms, and the trigonometric functions. These are the **elementary functions**. This lecture is a tour of the list, noting for each type its natural domain, its range, and its basic shape — information we will need constantly once we start differentiating and integrating.

5.1 Polynomials

Definition 5.1.

A **polynomial** is a function of the form

$$p(x) = a_0 + a_1x + a_2x^2 + \cdots + a_nx^n,$$

with real coefficients $a_0, a_1, \dots, a_n$. If $a_n \neq 0$, the integer n is the **degree** of p .

The lowest degrees have familiar names:

degree	form	name
0	$y = a_0$	constant
1	$y = a_0 + a_1x$	linear
2	$y = a_0 + a_1x + a_2x^2$	quadratic
3	$y = a_0 + a_1x + a_2x^2 + a_3x^3$	cubic

A polynomial is a finite sum of constant multiples of the **power functions** $1, x, x^2, x^3, \dots$, each of which is defined for every real number. Hence every polynomial has natural domain

$$\text{dom}(p) = \mathbb{R}$$

(unless we deliberately restrict it).

Remark 5.1. The pure powers x^n share a few features worth remembering. Every graph passes through $(0, 0)$ and $(1, 1)$. On the interval $[0, 1]$ a higher power lies *below* a lower one ($x^3 \leq x^2 \leq x$), while on $[1, \infty)$ the order reverses. The power x^n is even when n is even and odd when n is odd; for odd n it is strictly increasing on $\mathbb{R}$ (hence injective), while for even n it is not injective on $\mathbb{R}$.

5.2 Rational Functions

Definition 5.2.

A **rational function** is a quotient of two polynomials,

$$r(x) = \frac{p(x)}{q(x)}, \quad q \not\equiv 0.$$

Its natural domain is $\mathbb{R} \setminus Z$, where

$$Z = \{x \in \mathbb{R} : q(x) = 0\}$$

is the (finite) set of zeros of the denominator.

Example 5.1.

- $\frac{x}{x^2 + 1}$ is a rational function. Since $x^2 + 1 > 0$ for every real x , the set Z is empty and the domain is all of $\mathbb{R}$.
- $\frac{\sqrt{x}}{x^3 - x}$ is **not** a rational function: $\sqrt{x}$ is not a polynomial. (It is, however, still an elementary function, as it is the arithmetic combination of two elementary functions - see next section)

5.3 Algebraic and Power Functions

Definition 5.3.

A **power function** is $f(x) = x^\alpha$ with a fixed real exponent $\alpha \in \mathbb{R}$. For $x > 0$ this is defined for every α (via $x^\alpha = \exp(\alpha \ln x)$, once we have $\exp$ and $\ln$). For an **integer** exponent the domain extends to

$$\text{dom}(x^\alpha) = \begin{cases} \mathbb{R} & \alpha \geq 0, \\ \mathbb{R} \setminus \{0\} & \alpha < 0. \end{cases}$$

Remark 5.2. For a non-integer exponent one usually keeps the domain $[0, \infty)$ (or $(0, \infty)$ if $\alpha < 0$). The roots $x^{1/2} = \sqrt{x}$, $x^{1/3} = \sqrt[3]{x}$, ... are examples; odd roots such as $\sqrt[3]{x}$ can in fact be extended to all of $\mathbb{R}$. Power functions with rational exponent, and more generally functions obtained from polynomials by the four arithmetic operations together with root extraction, are called **algebraic functions**. The exponential and logarithm below are *not* algebraic — they are **transcendental**.

5.4 Exponential Functions

Definition 5.4.

Fix a **base** $a > 0$. The **exponential function** with base a ,

$$\exp_a : \mathbb{R} \rightarrow \mathbb{R}, \quad \exp_a(x) = a^x,$$

is defined for every real x . For $a \neq 1$ it takes only positive values, and every positive value is attained:

$$\text{dom}(\exp_a) = \mathbb{R}, \quad \text{ran}(\exp_a) = (0, \infty) = \mathbb{R}^{>0}.$$

For $a > 1$ the function a^x is strictly increasing; for $0 < a < 1$ it is strictly decreasing; for $a = 1$ it is the constant 1. In all cases $a^0 = 1$ and $a^{x+y} = a^x a^y$.

5.5 Logarithms

Definition 5.5.

Fix a base $a > 0$ with $a \neq 1$. Then $\exp_a$ is injective from $\mathbb{R}$ onto $(0, \infty)$, and its inverse function (in the sense of Definition 4.8) is the **logarithm to base a** ,

$$\log_a : (0, \infty) \rightarrow \mathbb{R}, \quad \log_a(a^x) = x, \quad a^{\log_a y} = y.$$

Thus

$$\text{dom}(\log_a) = (0, \infty), \quad \text{ran}(\log_a) = \mathbb{R}.$$

Because it is the inverse of a^x , the graph of $\log_a$ is the reflection of the graph of a^x across the line $y = x$ (Remark 4.3). It converts products to sums: $\log_a(xy) = \log_a x + \log_a y$.

5.6 Trigonometric Functions

Let x be a real number, thought of as an angle measured in **radians**. Rotating the point $(1, 0)$ counter-clockwise through the angle x about the origin lands on a point of the unit circle whose coordinates are, by definition, $(\cos x, \sin x)$.

Both $\sin$ and $\cos$ are defined for all $x \in \mathbb{R}$, take values in $[-1, 1]$, and repeat after a full turn.

Definition 5.6.

A function f is **periodic with period** $P > 0$ if $f(x + P) = f(x)$ for every x in its domain.

Thus $\sin$ and $\cos$ are periodic with period 2π . The **tangent**

$$\tan x = \frac{\sin x}{\cos x}$$

is defined wherever $\cos x \neq 0$, i.e. for $x \neq \frac{\pi}{2} + k\pi$ with $k \in \mathbb{Z}$, and is periodic with period π .

Remark 5.3. A general **sinusoid** has the form

$$f(x) = A \sin(bx) + C,$$

with **amplitude** $|A|$ (half the vertical spread of the graph), **period** $\frac{2\pi}{b}$, and **vertical offset** C (the horizontal line the graph oscillates around).

Exercise 5.1. The function $f(x) = 3 \sin(2x) + 1$ is a sinusoid. Find its amplitude, period, and range, and the smallest $x > 0$ at which it attains its maximum.

5.7 Summary

- **Polynomials** (Definition 5.1) are finite combinations of the powers $1, x, x^2, \dots$; they are defined on all of $\mathbb{R}$.
- **Rational functions** (Definition 5.2) are quotients of polynomials, defined off the finite zero set of the denominator.
- **Power functions** x^α (Definition 5.3) have domain $\mathbb{R}$ or $\mathbb{R} \setminus \{0\}$ for integer exponents, and are typically taken on $[0, \infty)$ otherwise; together with roots they form the **algebraic** functions.
- **Exponentials** a^x (Definition 5.4) have domain $\mathbb{R}$ and range $(0, \infty)$; their inverses, the **logarithms** $\log_a$ (Definition 5.5), have domain $(0, \infty)$ and range $\mathbb{R}$.

- The **trigonometric functions** $\sin, \cos$ are 2π -periodic (Definition [5.6](#)) with range $[-1, 1]$; $\tan = \sin / \cos$ is π -periodic and undefined where $\cos$ vanishes.

6 Continuity

Companion reading

This lecture corresponds to C&J §1.2.d.

Intuitively, a function is continuous if a small change in the independent variable produces only a small change in the dependent variable — the graph has no jumps, gaps, or blow-ups. Before making this precise, it helps to see why the intuitive idea is not yet a *definition*: it uses the word “small” twice without saying how small, or how the two smallnesses are related.

Example 6.1.

Let

$$f(x) = \begin{cases} 1 & x \geq 0 \\ 0 & x < 0. \end{cases}$$

At $x_0 = 0$, an arbitrarily small change in x — moving from $x = 0$ to some $x < 0$ — causes f to jump from 1 to 0, a change of a full unit. No matter how small a step we take to the left of 0, the output change stays large. We call this a **jump discontinuity** at $x_0 = 0$.

Q: How do we make “small change forces small change” precise? The idea is to formalize “small” using neighborhoods.

Definition 6.1.

Given $a \in \mathbb{R}$ and $\varepsilon > 0$, the ε -**neighborhood** of a is the interval

$$(a - \varepsilon, a + \varepsilon).$$

Think of ε as a *tolerance*: saying “ x is within tolerance of a ” means $x \in (a - \varepsilon, a + \varepsilon)$, i.e. $|x - a| < \varepsilon$. “Small change” becomes “ ε small.”

Fix a point x_0 . We want: the distance between $f(x)$ and $f(x_0)$ stays within tolerance ε , i.e.

$$|f(x) - f(x_0)| < \varepsilon,$$

whenever x is *sufficiently close* to x_0 — say, within some tolerance δ of x_0 . The key point is that this has to work for **any** $\varepsilon > 0$, no matter how small; but we are allowed to choose δ *in response to* ε — a smaller tolerance on the output may well force a smaller tolerance on the input.

6.1 The ε - δ Definition

Definition 6.2.

Let $f: D \rightarrow \mathbb{R}$ with $D \subseteq \mathbb{R}$, and let $x_0 \in D$. We say f is **continuous at** x_0 if for every $\varepsilon > 0$ there exists $\delta > 0$ such that

$$x \in D \text{ and } |x - x_0| < \delta \quad \Rightarrow \quad |f(x) - f(x_0)| < \varepsilon.$$

Example 6.2.

Take f as in Example 6.1 and $x_0 = 0$. Pick $\varepsilon = \frac{1}{2}$. No matter how small we choose $\delta > 0$, there is always some $x < 0$ with $|x - x_0| < \delta$ (i.e. x is δ -close to $x_0 = 0$), and for any such x ,

$$|f(x) - f(x_0)| = |0 - 1| = 1 > \frac{1}{2}.$$

So no δ works for this ε , and f is **not** continuous at $x_0 = 0$.

Remark 6.1. Continuity is a property of f **at a point**. The function f above is continuous at every $x_0 \neq 0$ (Exercise 6.1) and discontinuous only at $x_0 = 0$ — a function can be continuous at some points of its domain and discontinuous at others.

Exercise 6.1. Show that f from Example 6.1 is continuous at every $x_0 \neq 0$.

6.1.1 The ε - δ game

Remark 6.2. It helps to think of the definition as a game between an opponent and you. The opponent challenges you with some $\varepsilon > 0$; you must exhibit a matching $\delta > 0$ so that

$$|x - x_0| < \delta \quad \Rightarrow \quad |f(x) - f(x_0)| < \varepsilon.$$

f is continuous at x_0 exactly when you can win this game **no matter which** ε **the opponent picks**.

Example 6.3.

$g(x) = 3x + 2$ is continuous at every $x_0 \in \mathbb{R}$.

Proof. Let $\varepsilon > 0$ and $x_0 \in \mathbb{R}$ be arbitrary. We want $|g(x) - g(x_0)| < \varepsilon$. Compute

$$|g(x) - g(x_0)| = |(3x + 2) - (3x_0 + 2)| = |3x - 3x_0| = 3|x - x_0|.$$

So if we let $\delta = \varepsilon/3$, then

$$|x - x_0| < \delta \quad \Rightarrow \quad |g(x) - g(x_0)| = 3|x - x_0| < 3\delta = \varepsilon.$$

□

Exercise 6.2. Show, directly from Definition 6.2, that $h(x) = \frac{1}{2}x - 1$ is continuous at every $x_0 \in \mathbb{R}$. (What choice of δ in terms of ε makes the argument work?)

A Jupyter demonstration accompanying this lecture lets you play the ε - δ game interactively: pick x_0 and a δ , and watch the resulting range of f -values shrink (or fail to shrink) as $\delta \rightarrow 0$.

6.2 Types of Discontinuities

If f fails to be continuous at x_0 , the failure typically has one of three flavors.

Remark 6.3. Jump discontinuities. The values of f approach one number from the left of x_0 and a different number from the right, as in Example 6.1. We will make this precise next week with *one-sided limits*: a jump discontinuity is exactly a point where the one-sided limits exist but disagree.

Remark 6.4. Unbounded (infinite) discontinuities. Consider

$$f(x) = \begin{cases} 1/x^2 & x \neq 0 \\ 0 & x = 0. \end{cases}$$

As $x \rightarrow 0$, $f(x) \rightarrow \infty$, so $|f(x) - f(0)|$ is **unbounded** on every neighborhood of 0 — no ε -tolerance, however large, can be met near $x_0 = 0$ by any δ .

Remark 6.5. Oscillation discontinuities. Consider

$$g(x) = \begin{cases} \sin(1/x) & x \neq 0 \\ 0 & x = 0. \end{cases}$$

Near $x_0 = 0$, g oscillates wildly: arbitrarily close to 0 we can always find points where g equals 1 and points where g equals -1 . Indeed, for any $n \in \mathbb{N}$, taking

$$x_1 = \frac{1}{2\pi n + \pi/2} \quad \text{gives} \quad \sin(1/x_1) = 1,$$

and similarly $x_{-1} = \frac{1}{2\pi n + 3\pi/2}$ gives $\sin(1/x_{-1}) = -1$; both $x_1, x_{-1} \rightarrow 0$ as $n \rightarrow \infty$. So on every neighborhood of 0, no matter how small, g takes values ε apart for ε close to 2 — g cannot be made continuous at 0 by any choice of $g(0)$.

Example 6.4.

Contrast this with

$$x \sin(1/x) \quad (x \neq 0), \quad \text{value } 0 \text{ at } x = 0,$$

which **is** continuous at $x_0 = 0$: although $\sin(1/x)$ still oscillates without settling down, it is squeezed between -1 and 1 , so $|x \sin(1/x) - 0| \leq |x| \rightarrow 0$ as $x \rightarrow 0$. Oscillation alone does not cause a discontinuity — it must fail to be *damped out* as $x \rightarrow x_0$.

6.3 Summary

- The intuitive idea “small change in, small change out” is formalized using **ε -neighborhoods** (Definition 6.1): a tolerance interval $(a - \varepsilon, a + \varepsilon)$ around a point.
- f is **continuous at** x_0 (Definition 6.2) if for every tolerance $\varepsilon > 0$ on the output, there is a matching tolerance $\delta > 0$ on the input, so that x within δ of x_0 forces $f(x)$ within ε of $f(x_0)$ — think of it as a game (Remark 6.2) you must be able to win against any ε .
- Continuity is a property of f **at a point** (Remark 6.1); a function can be continuous at some points and not others.
- Failures of continuity come in three standard flavors: **jump** discontinuities (Remark 6.3), where one-sided behavior disagrees; **unbounded** discontinuities (Remark 6.4), where f blows up; and **oscillation** discontinuities (Remark 6.5), where f never settles down near x_0 — unless the oscillation is damped, as in $x \sin(1/x)$ (Example 6.4), which remains continuous.

7 Activity: Properties of Functions

About this activity

An in-class, pencil-and-paper worksheet. It revisits the properties of functions built up in the previous chapters — *injective*, *monotone*, *even/odd*, and, above all, *continuous* — by working them out in full for two specific functions. Companion reading: C&J §1.2.d.

7.1 Background

We have met several properties a function can have: **injective**, **monotone**, **even or odd**, and — most important — **continuous**. In this activity we investigate all of them for the two functions

$$f(x) = \begin{cases} \frac{x}{|x|} & x \neq 0, \\ 0 & x = 0, \end{cases} \quad g(x) = \begin{cases} \frac{1}{x^2} & x \neq 0, \\ 0 & x = 0. \end{cases}$$

7.2 Step 1: Draw the graphs

Sketch the graph of each function on its own set of axes. Plot a few sample points first (e.g. $x = \pm\frac{1}{2}, \pm 1, \pm 2, \pm 3$), then connect them.

7.3 Step 2: Make an educated guess

From your graphs, try to answer the following.

1. Is the function injective?
2. Is it monotone? If so, is it *strictly* monotone? If not, can you find intervals — as large as possible — on which it is (strictly) monotone?
3. Is it even, odd, or neither?




Figure 7.1: Blank axes for sketching f (left) and g (right).

4. Is it continuous everywhere? If not, at which x is it discontinuous?

Record short answers in the table; you will verify each one in the steps that follow.

Property	f	g
Injective? (yes / no)		
Monotone? Strictly? On which interval(s)?		
Even, odd, or neither?		
Continuous everywhere? If not, at which x ?		

7.4 Verification: injectivity

Neither f nor g is injective. To prove this, it is enough to exhibit two inputs $x \neq y$ with the same output.

Exercise 7.1. Find explicit values:

- for f : $x = \underline{\hspace{1cm}}$ and $y = \underline{\hspace{1cm}}$ with $f(x) = f(y)$;
- for g : $x = \underline{\hspace{1cm}}$ and $y = \underline{\hspace{1cm}}$ with $g(x) = g(y)$.

7.5 Verification: even / odd

Exercise 7.2. For whichever symmetry you guessed, verify it directly from the definition of the function: compute $f(-x)$ and check that it equals $-f(x)$ (odd) or $f(x)$ (even). Then do the same for g .

7.6 Verification: monotonicity

Exercise 7.3. Verify your monotonicity guess on a suitable interval, directly from the definition. For instance, to show a function is increasing on an interval, assume $x \leq y$ in that interval and deduce $f(x) \leq f(y)$ (use strict inequalities for *strictly* increasing); argue similarly for decreasing.

7.7 Continuity: the ε - δ game

Both f and g are continuous at $x_0 = 1$. Play the ε - δ game from Definition 6.2.

Exercise 7.4.

1. Take $\varepsilon = 1$. Find a number $\delta > 0$ such that

$$|x - 1| < \delta \quad \Rightarrow \quad |f(x) - f(1)| < 1$$

(recall $f(1) = 1$). Then find such a δ for g .

2. Repeat with a smaller tolerance, $\varepsilon = \frac{1}{10}$. Does the argument really change?
3. Can you produce a δ that works for an **arbitrary** $\varepsilon > 0$?

7.8 Discontinuities

Both functions are discontinuous at $x_0 = 0$. Here it is *you* who loses the game.

Exercise 7.5. Find an $\varepsilon > 0$ for which **no** δ works — that is, however close x is taken to 0, there is always some such x with

$$|f(x) - f(0)| \geq \varepsilon.$$

Give such an ε for f , and one for g .

Part III

Limits

8 Sequences and Convergence

8.1 Sequences

A **sequence** is an unending list of real numbers

$$a_1, a_2, a_3, \dots$$

We write $(a_n)_{n \in \mathbb{N}}$, or just (a_n) , and call a_n the **n -th term**. Formally, a sequence is nothing but a function

$$a: \mathbb{N} \rightarrow \mathbb{R},$$

where we write a_n instead of $a(n)$.

Example 8.1.

$$\begin{aligned} a_n &= \frac{1}{n} : & 1, \frac{1}{2}, \frac{1}{3}, \frac{1}{4}, \dots \\ b_n &= \frac{(-1)^{n+1}}{n} : & 1, -\frac{1}{2}, \frac{1}{3}, -\frac{1}{4}, \dots \\ c_n &= \frac{n}{n+1} : & \frac{1}{2}, \frac{2}{3}, \frac{3}{4}, \frac{4}{5}, \dots \\ d_n &= \sqrt[n]{p}, \quad p > 0 \text{ fixed} \\ e_n &= \alpha^n, \quad \alpha > 1. \end{aligned}$$

8.2 The Limit of a Sequence

Look at $a_n = \frac{1}{n}$. What happens to the terms as n gets larger and larger? The values get smaller and smaller, coming **arbitrarily close to** 0 — though no term is ever equal to 0. We want to make “arbitrarily close” precise.

The idea is to use ε -neighborhoods (Definition 6.1): given a tolerance $\varepsilon > 0$, we ask that *from some point on*, every term of the sequence lies within ε of 0.

Definition 8.1.

We say (a_n) **converges to** 0, written

$$\lim_{n \rightarrow \infty} a_n = 0 \quad \text{or} \quad a_n \rightarrow 0,$$

if for every $\varepsilon > 0$ there exists $N \in \mathbb{N}$ such that

$$|a_n| = |a_n - 0| < \varepsilon \quad \text{for all } n \geq N.$$

The number N is allowed to depend on ε — a smaller tolerance will in general require going further out in the sequence. To emphasize this dependence one sometimes writes $N(\varepsilon)$.

Example 8.2.

$$\lim_{n \rightarrow \infty} \frac{1}{n} = 0.$$

Proof. Let $\varepsilon > 0$. We must find N so that $\frac{1}{n} < \varepsilon$ for all $n \geq N$. Now

$$\frac{1}{n} < \varepsilon \quad \Leftrightarrow \quad \frac{1}{\varepsilon} < n,$$

so any $N \in \mathbb{N}$ with $N > \frac{1}{\varepsilon}$ works: if $n \geq N > \frac{1}{\varepsilon}$, then $\frac{1}{n} < \varepsilon$.

Concretely, for $\varepsilon = 0.05 = \frac{5}{100}$ we need $n > \frac{100}{5} = 20$, so $N = 21$ will do. $\square$

Exercise 8.1. Show that $\lim_{n \rightarrow \infty} \frac{(-1)^{n+1}}{n} = 0$. (Note $|b_n| = \frac{1}{n}$.)

8.3 Convergence to a General Limit

Now consider $c_n = \frac{n}{n+1}$. The terms get closer and closer to 1, which suggests $\lim_{n \rightarrow \infty} c_n = 1$. The definition is the same as before, with 0 replaced by the candidate limit.

Definition 8.2.

Let $a \in \mathbb{R}$. We say (c_n) **converges to** a , written $\lim_{n \rightarrow \infty} c_n = a$ or $c_n \rightarrow a$, if for every $\varepsilon > 0$ there exists $N \in \mathbb{N}$ such that

$$|c_n - a| < \varepsilon \quad \text{for all } n \geq N.$$

Equivalently: from the N -th term on, all terms lie in the ε -neighborhood $(a - \varepsilon, a + \varepsilon)$. A sequence is **convergent** if it converges to some a , and **divergent** otherwise.

Example 8.3.

$\lim_{n \rightarrow \infty} \frac{n}{n+1} = 1$. Here

$$|c_n - 1| = \left| \frac{n}{n+1} - 1 \right| = 1 - \frac{n}{n+1} = \frac{1}{n+1}.$$

Let us check two tolerances. For $\varepsilon = 1$ we need $\frac{1}{n+1} < 1$, which holds for every $n \geq 1$; so $N = 1$ works. For $\varepsilon = \frac{1}{4}$ we need

$$\frac{1}{n+1} < \frac{1}{4} \quad \Leftrightarrow \quad 3 < n,$$

so $N = 4$ works.

In general, $|c_n - 1| = \frac{1}{n+1} < \varepsilon \quad \Leftrightarrow \quad n > \frac{1}{\varepsilon} - 1$, so any $N > \frac{1}{\varepsilon} - 1$ works, and $c_n \rightarrow 1$.

Remark 8.1. When a limit exists it is **unique**: if $c_n \rightarrow a$ and $c_n \rightarrow a'$, then for every $\varepsilon > 0$ we have $|a - a'| \leq |a - c_n| + |c_n - a'| < 2\varepsilon$ for large n , so $|a - a'| = 0$. This is what justifies speaking of *the* limit and the notation $\lim_{n \rightarrow \infty} c_n$.

8.3.1 Divergence

To show a sequence **diverges** we must show it converges to no limit at all. Negating Definition 8.2: (c_n) does not converge to a if there is some $\varepsilon > 0$ such that, no matter how far out we go, some later term is at least ε away from a .

Example 8.4.

The sequence $((-1)^n) = -1, 1, -1, 1, \dots$ diverges.

Proof. Suppose $(-1)^n \rightarrow a$ for some a . Take $\varepsilon = 1$. Then $|(-1)^n - a| < 1$ for all large n ; applying this to an even and an odd n gives $|1 - a| < 1$ and $|-1 - a| < 1$. But then

$$2 = |1 - (-1)| \leq |1 - a| + |a - (-1)| < 1 + 1 = 2,$$

a contradiction. □

Exercise 8.2. Investigate the remaining sequences from Example 8.1: $d_n = \sqrt[n]{p}$ (with $p > 0$ fixed) and $e_n = \alpha^n$ (with $\alpha > 1$). One converges; the other diverges. Which is which, and what is the limit?

Remark 8.2. For reference: $\sqrt[n]{p} \rightarrow 1$ for every $p > 0$, while α^n diverges when $\alpha > 1$ (its terms grow without bound). Both facts are proved in C&J §1.6.

8.4 Limit Laws

Computing $N(\varepsilon)$ by hand for every sequence quickly becomes tedious. The following rules let us build limits of complicated sequences out of simpler ones.

Proposition 8.1 (Arithmetic of limits). *Suppose $a_n \rightarrow a$ and $b_n \rightarrow b$. Then*

$$\begin{aligned} a_n + b_n &\rightarrow a + b, & a_n - b_n &\rightarrow a - b, \\ a_n \cdot b_n &\rightarrow a \cdot b, & \frac{a_n}{b_n} &\rightarrow \frac{a}{b} \quad (\text{provided } b \neq 0). \end{aligned}$$

Example 8.5.

Find $\lim_{n \rightarrow \infty} \frac{2n^2 - 1}{n^2 + n + 1}$.

Divide numerator and denominator by n^2 :

$$\frac{2n^2 - 1}{n^2 + n + 1} = \frac{2 - \frac{1}{n^2}}{1 + \frac{1}{n} + \frac{1}{n^2}}.$$

Now $\frac{1}{n} \rightarrow 0$ and $\frac{1}{n^2} \rightarrow 0$, so by the limit laws the numerator tends to $2 - 0 = 2$ and the denominator to $1 + 0 + 0 = 1$. Since the denominator limit is nonzero,

$$\lim_{n \rightarrow \infty} \frac{2n^2 - 1}{n^2 + n + 1} = \frac{2}{1} = 2.$$

8.5 Comparison and the Squeeze Theorem

Exercise 8.3. Warm-up. Show that $\lim_{n \rightarrow \infty} \frac{1}{2^n} = 0$.

Rather than hunting for $N(\varepsilon)$ from scratch every time, it is often easier to compare a sequence to one whose limit we already know.

Proposition 8.2 (Comparison test). *Suppose $0 \leq b_n \leq a_n$ for all n , and $a_n \rightarrow 0$. Then $b_n \rightarrow 0$.*

Proof. Let $\varepsilon > 0$. Since $a_n \rightarrow 0$ there is $N \in \mathbb{N}$ with $a_n < \varepsilon$ for all $n \geq N$. Then for $n \geq N$,

$$0 \leq b_n \leq a_n < \varepsilon,$$

so $|b_n - 0| = b_n < \varepsilon$. Hence $b_n \rightarrow 0$. □

We can generalize this by sandwiching a sequence between two others that converge to the *same* limit works for any target value.

Proposition 8.3 (Squeeze theorem). *Suppose $a_n \leq b_n \leq c_n$ for all n , and $a_n \rightarrow L$, $c_n \rightarrow L$ for the same $L \in \mathbb{R}$. Then $b_n \rightarrow L$.*

Proof. Let $\varepsilon > 0$. Choose N_1 with $|a_n - L| < \varepsilon$ for $n \geq N_1$, and N_2 with $|c_n - L| < \varepsilon$ for $n \geq N_2$. Let $N = \max(N_1, N_2)$. For $n \geq N$,

$$L - \varepsilon < a_n \leq b_n \leq c_n < L + \varepsilon,$$

so $|b_n - L| < \varepsilon$. Hence $b_n \rightarrow L$. $\square$

Remark 8.3. Proposition 8.2 is the special case of the squeeze theorem where the lower bounding sequence is constantly 0 and $L = 0$.

8.6 Monotone and Bounded Sequences

Definition 8.3.

A sequence (a_n) is **increasing** if $a_n \leq a_{n+1}$ for all n , and **decreasing** if $a_n \geq a_{n+1}$ for all n . It is **monotone** if it is increasing or decreasing.

Definition 8.4.

A sequence (a_n) is **bounded** if there is a constant C such that $|a_n| \leq C$ for all n .

Theorem 8.1 (Monotone Convergence Theorem). *Every bounded monotone sequence converges.*

Remark 8.4. This theorem expresses the same completeness of $\mathbb{R}$ as the Nested Interval Axiom from [The Nested Interval Axiom](#) — each can be derived from the other. We will not reprove the general statement here; instead we isolate the paired form that we actually use below, which follows directly from the axiom together with the comparison test.

Corollary 8.1. *Suppose (a_n) is increasing, (b_n) is decreasing, $a_n \leq b_n$ for all n , and $b_n - a_n \rightarrow 0$. Then (a_n) and (b_n) converge to a common limit c , the unique point in $\bigcap_n [a_n, b_n]$.*

Proof. Since (a_n) is increasing and (b_n) is decreasing with $a_n \leq b_n$, the intervals $[a_n, b_n]$ are nested: $[a_{n+1}, b_{n+1}] \subseteq [a_n, b_n]$ for all n . By the Nested Interval Axiom there is a point c with $a_n \leq c \leq b_n$ for every n . Then

$$0 \leq c - a_n \leq b_n - a_n \rightarrow 0 \quad \text{and} \quad 0 \leq b_n - c \leq b_n - a_n \rightarrow 0,$$

so the comparison test (Proposition 8.2) gives $a_n \rightarrow c$ and $b_n \rightarrow c$. $\square$

8.7 Applications

Lemma 8.1 (Bernoulli's inequality). *For $c \geq -1$ and $n \in \mathbb{N}$,*

$$(1 + c)^n \geq 1 + nc.$$

(For $c \geq 0$ one can prove this by considering just the first two terms of the [binomial expansion](#) of $(1 + c)^n$. The remaining terms are all non-negative, so by leaving them out, we obtain something smaller.)

8.7.1 Roots: $\sqrt[n]{p} \rightarrow 1$

Theorem 8.2.

For every $p > 0$, $\lim_{n \rightarrow \infty} \sqrt[n]{p} = 1$.

Proof. For $p = 1$ the sequence is constant, so there is nothing to show.

Suppose $p > 1$. Write $a_n := \sqrt[n]{p} = 1 + \delta_n$; since $p > 1$ we have $a_n > 1$, so $\delta_n > 0$. Raising to the n -th power and applying Bernoulli's inequality (with $c = \delta_n \geq 0$),

$$p = (1 + \delta_n)^n \geq 1 + n\delta_n,$$

so $0 < \delta_n \leq \frac{p-1}{n}$. Since $\frac{p-1}{n} \rightarrow 0$, the comparison test gives $\delta_n \rightarrow 0$, and hence $a_n = 1 + \delta_n \rightarrow 1$.

Suppose $0 < p < 1$. Then $q := 1/p > 1$, and by the case just proved $\sqrt[n]{q} \rightarrow 1$. Since $\sqrt[n]{q} \neq 0$ for all n and its limit is nonzero, the quotient rule in [Proposition 8.1](#) gives

$$\sqrt[n]{p} = \frac{1}{\sqrt[n]{q}} \rightarrow \frac{1}{1} = 1.$$

□

This proves the claim recorded in [Remark 8.2](#) for $d_n = \sqrt[n]{p}$.

8.7.2 The Number e

Proposition 8.4.

The sequence $e_n := \left(1 + \frac{1}{n}\right)^n$ is increasing.

Proof. We must show

$$\left(1 + \frac{1}{n}\right)^n \leq \left(1 + \frac{1}{n+1}\right)^{n+1}. \quad (*)$$

Writing both sides with a common fractional base and splitting off one factor of $\frac{n+2}{n+1}$ on the right,

$$(*) \iff \left(\frac{n+1}{n}\right)^n \leq \left(\frac{n+2}{n+1}\right)^n \cdot \frac{n+2}{n+1} \iff \frac{n+1}{n+2} \leq \left(\frac{n(n+2)}{(n+1)^2}\right)^n.$$

Since $(n+1)^2 = n^2 + 2n + 1$, we have $n(n+2) = n^2 + 2n = (n+1)^2 - 1$, so

$$\left(\frac{n(n+2)}{(n+1)^2}\right)^n = \left(1 - \frac{1}{(n+1)^2}\right)^n.$$

By Bernoulli's inequality (with $c = -\frac{1}{(n+1)^2} \geq -1$),

$$\left(1 - \frac{1}{(n+1)^2}\right)^n \geq 1 - \frac{n}{(n+1)^2}.$$

So it suffices to show $\frac{n+1}{n+2} \leq 1 - \frac{n}{(n+1)^2}$, i.e.

$$\frac{1}{n+2} \geq \frac{n}{(n+1)^2} \iff (n+1)^2 \geq n(n+2) \iff n^2 + 2n + 1 \geq n^2 + 2n,$$

which is always true. Chaining the inequalities back together proves (*). $\square$

Exercise 8.4. Show that $f_n := \left(1 + \frac{1}{n}\right)^{n+1}$ is decreasing. (Adapt the computation in the proof of Proposition 8.4.)

Theorem 8.3.

The limit

$$e := \lim_{n \rightarrow \infty} \left(1 + \frac{1}{n}\right)^n$$

exists; numerically $e \approx 2.71828 \dots$

Proof. Let $e_n = (1 + \frac{1}{n})^n$ and $f_n = (1 + \frac{1}{n})^{n+1} = e_n (1 + \frac{1}{n})$, so $e_n \leq f_n$ for all n . By Proposition 8.4, (e_n) is increasing; by Exercise 8.4, (f_n) is decreasing. In particular $e_n \leq f_n \leq f_1 = 4$ for all n , so (e_n) is bounded.

Moreover

$$0 \leq f_n - e_n = e_n \cdot \frac{1}{n} \leq \frac{f_1}{n} = \frac{4}{n} \rightarrow 0,$$

so $f_n - e_n \rightarrow 0$ by the comparison test. By Corollary 8.1, (e_n) and (f_n) converge to a common limit — which we call e . Numerically, $e_{10} \approx 2.594$, $e_{100} \approx 2.705$, $e_{1000} \approx 2.7169$, consistent with $e \approx 2.71828 \dots$ $\square$

8.8 Summary

- A **sequence** is a function $a: \mathbb{N} \rightarrow \mathbb{R}$, written (a_n) .
- (a_n) **converges to** a (Definition 8.2) if for every tolerance $\varepsilon > 0$ there is an index N such that $|a_n - a| < \varepsilon$ for all $n \geq N$ — from some point on, all terms lie in the ε -neighborhood of a . The special case $a = 0$ is Definition 8.1.
- The limit, when it exists, is **unique** (Remark 8.1). A sequence with no limit, such as $((-1)^n)$ (Example 8.4), is **divergent**.
- The **limit laws** (Proposition 8.1) compute limits of sums, differences, products, and quotients termwise; dividing through by the highest power of n turns a ratio of polynomials into a form where they apply (Example 8.5).
- The **comparison test** and **squeeze theorem** (Proposition 8.2, Proposition 8.3) let us find a limit by trapping a sequence between two others whose limits we already know.
- A **bounded monotone sequence converges** (Theorem 8.1), a fact equivalent to the Nested Interval Axiom; the paired form (Corollary 8.1) — an increasing and a decreasing sequence squeezing together — is the version we use in practice.
- **Bernoulli's inequality** $(1 + c)^n \geq 1 + nc$ (Lemma 8.1) drives two applications: every $\sqrt[n]{p} \rightarrow 1$ (Theorem 8.2), and $(1 + \frac{1}{n})^n$ increases while $(1 + \frac{1}{n})^{n+1}$ decreases to the same limit, defining the number $e \approx 2.71828 \dots$ (Theorem 8.3).

9 Convergence of Sequences: Reading Data

```
#
# !!! RUN THIS FIRST !!!
#
# (loads the data and defines all necessary functions)
#

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

pd.set_option('display.max_rows', 20) # long tables show head/tail with '...' automatic

df = pd.read_csv('data/rc_circuit_data.csv')

def a(n):
    """Return the n-th term of the sequence, a_n, as a plain float."""
    row = df.loc[df['n'] == n]
    if row.empty:
        raise ValueError(f"n = {n} is out of range (data covers n = 0 .. {df['n'].max()})")
    return float(row['measured_value'].values[0])

def find_smallest_N(data, L_guess, epsilon, value_col='measured_value'):
    """Smallest N (present in the data) such that all later points are within epsilon of
    vals = data[value_col].values
    n_vals = data['n'].values
    for i in range(len(vals)):
        if np.all(np.abs(vals[i:] - L_guess) < epsilon):
            return int(n_vals[i])
    return None
```

```

def check_convergence(data, L_guess, N_guess, epsilon, value_col='measured_value', reveal=True):
    vals = data[value_col].values
    n_vals = data['n'].values
    mask = n_vals >= N_guess

    if not mask.any():
        print(f"N_guess = {N_guess} is beyond the range of available data (max n = {n_vals.max()})")
        return False

    tail_n = n_vals[mask]
    tail_vals = vals[mask]
    bad = np.abs(tail_vals - L_guess) >= epsilon

    if bad.any():
        first_bad_n = tail_n[bad][0]
        print(f" Not valid: at n = {first_bad_n}, |a_n - {L_guess}| >= {epsilon}.")
        print(" The sequence leaves the epsilon-band after your N_guess. Try a larger N_guess "
              "or reconsider your L_guess.")
        return False

    print(f" Valid: for all n >= {N_guess}, |a_n - {L_guess}| < {epsilon}.")

    # Don't just hand over the smallest N -- only reveal it once the guess is close.
    best_N = find_smallest_N(data, L_guess, epsilon, value_col)
    if best_N is not None:
        gap = N_guess - best_N
        if gap == 0:
            print(" This is also the smallest N that works for this L_guess.")
        elif gap <= reveal_threshold:
            print(f" Close! The smallest N that works is N = {best_N} (you were only {gap} away)")
        else:
            print(" This N works, but it's probably not the smallest one -- "
                  "try a smaller N and see if it still holds.")
    return True

def plot_challenge(data, L_guess, N_guess=None, epsilon=None, value_col='measured_value',
                  zoom=True, buffer=50):
    """
    Plot the sequence with the guessed L (and epsilon band, and N_guess) marked.

    By default (zoom=True), the plot starts `buffer` samples before N_guess and

```

```

runs to the end of the data -- enough to see whether the sequence was
outside the band right before N_guess, and to see it stay inside the band
all the way to the end. Set zoom=False to see the whole sequence instead,
or pass a bigger `buffer` to look further back.
"""

plot_data = data
if zoom and N_guess is not None:
    start = max(int(data['n'].min()), N_guess - buffer)
    mask = data['n'] >= start
    plot_data = data[mask]

plt.figure(figsize=(8, 4))
plt.plot(plot_data['n'], plot_data[value_col], '.', alpha=0.6)
plt.axhline(L_guess, color='green', linestyle='--', label=f'L_guess = {L_guess}')
if epsilon is not None:
    plt.fill_between(plot_data['n'], L_guess - epsilon, L_guess + epsilon,
                     color='green', alpha=0.15, label=f'epsilon band (+/- {epsilon})')
if N_guess is not None:
    plt.axvline(N_guess, color='red', linestyle=':', label=f'N_guess = {N_guess}')
plt.xlabel('n')
plt.ylabel(value_col)
plt.legend()
plt.show()

```

You've been handed a log file, `data/rc_circuit_data.csv`, from a lab sensor. It records a voltage measurement `measured_value` at each sample index `n` (and the corresponding time `t_seconds`).

You do not know what circuit or process produced this data. Your job is to:

1. Explore the data and form a hypothesis about its limiting behavior.
2. Given a challenge value ε , **guess** the limit L and the index N after which the data stays within ε of L — then verify your guess rigorously, not just by eye.

Recall the definition:

$$\lim_{n \rightarrow \infty} a_n = L \iff \text{for all } \varepsilon > 0, \text{ there is } N \text{ s.t. } n \geq N \implies |a_n - L| < \varepsilon.$$

9.1 A Simple Query Function

Sometimes you just want a single term of the sequence, a_n , without pulling up the whole table. `a(n)` does that.

```
# try it:  
  
a(0), a(5000), a(9999)
```

9.2 Task 1: Explore the Data

The full log has 10,000 samples. Set `n_start` and `n_end` below to pick a window and re-run the cell to see the table and plot over just that range — e.g. zoom into the early samples to see the initial behavior, or into the tail to see whether it's still moving. Change the numbers and re-run as many times as you like.

```
# EDIT HERE  
  
n_start = 3000  
n_end = 9000  
  
#  
# No need to change anything down here  
#  
view = df[(df['n'] >= n_start) & (df['n'] <= n_end)]  
display(view)  
  
plt.figure(figsize=(8, 4))  
plt.plot(view['n'], view['measured_value'], '.', alpha=0.7)  
plt.xlabel('n')  
plt.ylabel('measured_value')  
plt.title(f'Sensor log, n = {n_start} to {n_end}')  
plt.show()
```

9.2.1 Questions to answer before moving on

- What value does the sequence appear to be approaching? Write down your best guess for L .

- Roughly how many samples does it take before the data “settles down” near that value?
- Does the data ever exceed your guessed L , or only approach it from below?
- Try $a(n)$ for a few large n (e.g. 7000, 8500, 9999) — does it match what the plot suggested?

9.3 Task 2: The ε Challenge

Your instructor assigns a value of ε . Using **only the data** (no formulas, no ground truth), you must:

1. State your guessed limit `L_guess`.
2. State the smallest index `N_guess` you believe works.
3. Check your answer with `check_convergence` below, which verifies directly against the definition: it checks whether **every** remaining data point from `N_guess` onward is within `epsilon` of your `L_guess`.

Note this checks self-consistency with your own guessed L — it does not tell you the “true” limit, because the whole point is that you don’t get to see that. Two different reasonable guesses for L can each have a valid N .

The checker also won’t just hand you the smallest possible N : if your guess is close (within 50), it’ll tell you the exact smallest N ; if you’re further off, it’ll just tell you to try smaller.

`plot_challenge` zooms in by default: it starts a bit before your `N_guess` and runs to the end of the data, so you can see both whether the sequence was outside the band just before `N_guess` and that it stays inside all the way to the end. Pass `zoom=False` if you want to see the whole sequence instead.

```
# --- Your answer ---
EPSILON = 0.08          # <- instructor-assigned challenge value

L_guess = 5             # TODO: your guessed limit
N_guess = 2400          # TODO: your guessed N

check_convergence(df, L_guess, N_guess, EPSILON)
plot_challenge(df, L_guess, N_guess, EPSILON)
```

9.4 Dataset 2: Damped Oscillator

This time there's no CSV and no noise. Instead, you're given the sequence directly as a **function** you can call for any index **n** — exactly the way a sequence (b_n) is usually presented in a calculus course. It models the position of an underdamped mechanical system (e.g. a suspension, a sensor needle) responding to a step input:

$$b_n = x_{ss} \left(1 - e^{-\zeta \omega_n t_n} \left(\cos(\omega_d t_n) + \frac{\zeta}{\sqrt{1 - \zeta^2}} \sin(\omega_d t_n) \right) \right), \quad t_n = n \cdot dt$$

with $\omega_d = \omega_n \sqrt{1 - \zeta^2}$. Unlike Dataset 1, this sequence **oscillates** around its limit before settling — so “getting close once” is not enough; you have to check it *stays* close.

Since it's an exact formula (not a fixed table), you can evaluate it at *any* **n** you like, no matter how large — there's no upper limit like there was with **a(n)**.

```
x_ss = 2.0      # steady-state value
zeta = 0.1      # damping ratio
wn = 10.0      # natural frequency (rad/s)
wd = wn * np.sqrt(1 - zeta**2)
dt2 = 0.001    # seconds per sample

def b(n):
    """Return the n-th term of the damped step-response sequence, b_n, as a
    plain float (or array of floats if n is array-like). No noise: this is
    the exact closed-form value."""
    scalar_input = np.ndim(n) == 0
    t = np.atleast_1d(np.asarray(n, dtype=float)) * dt2
    vals = x_ss * (1 - np.exp(-zeta * wn * t) *
                  (np.cos(wd * t) + (zeta / np.sqrt(1 - zeta**2)) * np.sin(wd * t)))
    return float(vals[0]) if scalar_input else vals

# try it -- note n can be far larger than anything in Dataset 1's table:
b(0), b(3000), b(1_000_000)
```

9.5 Task 3: Explore the Sequence

Same idea as Task 1: set **n_start** and **n_end** and re-run to see a table and plot of b_n over that window. Try zooming into the first couple thousand terms to see the

oscillation, then look further out to see it settle.

```
n_start = 0
n_end = 5000

n_vals = np.arange(n_start, n_end + 1)
view2 = pd.DataFrame({'n': n_vals, 'b_n': b(n_vals)})
display(view2)

plt.figure(figsize=(8, 4))
plt.plot(view2['n'], view2['b_n'], '.', alpha=0.7)
plt.xlabel('n')
plt.ylabel('b_n')
plt.title(f'Damped step response, n = {n_start} to {n_end}')
plt.show()
```

9.5.1 Questions to answer before moving on

- What value does b_n appear to be converging to? Compare to your guess for Dataset 1 — how did you decide this time, given there’s no noise to average out?
- Around how many terms does it take before the oscillation dies down to something you’d call “close” to the limit?
- Pick an n where the plot looks close to the limit, then check a later n — does it ever swing back out? Why might that matter for choosing N ?

9.6 Task 4: The ε Challenge, Again

Same process as Task 2 — guess L_{guess} and N_{guess} for a given epsilon , and check them. The `check_convergence`, `find_smallest_N`, and `plot_challenge` functions from Task 2 work unchanged here; we just need a table of b_n values to check against. Since b_n is exact (no noise), we can build that table out to a very large n — far past where anything interesting happens — so the check is effectively verifying “forever.”

```
n_check = np.arange(0, 20000)      # generous horizon -- well past where b_n settles
osc_df = pd.DataFrame({'n': n_check, 'b_n': b(n_check)})

# --- Your answer ---
EPSILON2 = 0.08                     # <- instructor-assigned challenge value
```

```

L_guess2 = 2          # TODO: your guessed limit
N_guess2 = 3000       # TODO: your guessed N

check_convergence(osc_df, L_guess2, N_guess2, EPSILON2, value_col='b_n')
plot_challenge(osc_df, L_guess2, N_guess2, EPSILON2, value_col='b_n')

```

9.7 Task 5: Convergence of $\sqrt[n]{p}$

Finally, we want to study the convergence behavior of $d_n = \sqrt[n]{p}$ for various p . Set the parameters in the cell below.

```

# Set p
p = 2      # Experiment with different p

# Set the range of n you want to consider
d_start = 1
d_end = 100

def d(n):
    return float(p**(1/n))

```

As above, to return just a single term of the sequence, use `d(n)`:

```
d(2)
```

Now let's plot the sequence:

```

d_array = np.arange(d_start, d_end)
d_vals = p**(1/d_array)
view3 = pd.DataFrame({'n': d_array, 'd_n': d_vals})
# display(view3)

plt.figure(figsize=(8, 4))
plt.plot(view3['n'], view3['d_n'], '.', alpha=0.7)
plt.xlabel('n')
plt.ylabel('d_n')
plt.title(f'Sequence n-th root of {p}')
plt.show()

```

Questions: - Does this sequence converge for all $p > 0$? - If so, does the limit depend on p ?